

New_York_TLC_1_EDA

April 1, 2025

1 New York TLC project

Part 1 - Explanatory Data Analysis

In this activity, you will examine data provided and prepare it for analysis. You will also design a professional data visualization that tells a story, and will help data-driven decisions for business needs.

The purpose of this project is to conduct exploratory data analysis on a provided data set. Your mission is to perform further EDA on this data with the aim of learning more about the variables.

The goal is to clean data set and create a visualization.

This activity has 4 parts:

Part 1: Imports, links, and loading

Part 2: Data Exploration (Data cleaning)

Part 3: Building visualizations

Part 4: Evaluate and share results

1. Identify any outliers:

- Use a boxplot to visualize the distribution of the data
- Use histograms to visualize the distribution of the data

2. How do you make the decision to keep or exclude outliers from any future models?

- There are three main options for dealing with outliers: keeping them as they are, deleting them, or reassigning them. Whether you keep outliers as they are, delete them, or reassign values is a decision that you make taking into account the nature of the outlying data and the assumptions of the model you are building. To help you make the decision, you can start with these general guidelines:
 - Delete them: If you are sure the outliers are mistakes, typos, or errors and the dataset will be used for modeling or machine learning, then you are more likely to decide to delete outliers. Of the three choices, you'll use this one the least.
 - Reassign them: If the dataset is small and/or the data will be used for modeling or machine learning, you are more likely to choose a path of deriving new values to replace the outlier values.
 - Leave them: For a dataset that you plan to do EDA/analysis on and nothing else, or for a dataset you are preparing for a model that is resistant to outliers, it is most likely that you are going to leave them in.

1.0.1 Task 1. Imports, links, and loading

For EDA of the data, import the data and packages that would be most helpful, such as pandas, numpy and matplotlib.

```
[11]: # Import packages and libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Note: As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[12]: # Load dataset into dataframe
df = pd.read_csv('2017_Yellow_Taxi_Trip_Data.csv')
```

1.0.2 Task 2a. Data exploration and cleaning

Decide which columns are applicable

The first step is to assess your data.

Given our scenario, which data columns are most applicable? Which data columns can I eliminate, knowing they won't solve our problem scenario?

Consider functions that help you understand and structure the data.

- head()
- describe()
- info()
- groupby()
- sortby()

What do the distributions of your variables tell you about the question you're asking or the problem you're trying to solve?

```
[13]: df.head()
```

```
[13]:
```

	Unnamed: 0	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	\
0	24870114	2	03/25/2017 8:55:43 AM	03/25/2017 9:09:47 AM	
1	35634249	1	04/11/2017 2:53:28 PM	04/11/2017 3:19:58 PM	
2	106203690	1	12/15/2017 7:26:56 AM	12/15/2017 7:34:08 AM	
3	38942136	2	05/07/2017 1:17:59 PM	05/07/2017 1:48:14 PM	
4	30841670	2	04/15/2017 11:32:20 PM	04/15/2017 11:49:03 PM	

	passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	\
0	6	3.34	1	N	
1	1	1.80	1	N	
2	1	1.00	1	N	

3	1	3.70	1	N
4	1	4.37	1	N

	PULocationID	DOLocationID	payment_type	fare_amount	extra	mta_tax	\
0	100	231	1	13.0	0.0	0.5	
1	186	43	1	16.0	0.0	0.5	
2	262	236	1	6.5	0.0	0.5	
3	188	97	1	20.5	0.0	0.5	
4	4	112	2	16.5	0.5	0.5	

	tip_amount	tolls_amount	improvement_surcharge	total_amount
0	2.76	0.0	0.3	16.56
1	4.00	0.0	0.3	20.80
2	1.45	0.0	0.3	8.75
3	6.39	0.0	0.3	27.69
4	0.00	0.0	0.3	17.80

```
[14]: df.shape
```

```
[14]: (22699, 18)
```

Use describe...

```
[15]: df.describe()
```

```
[15]:
```

	Unnamed: 0	VendorID	passenger_count	trip_distance	\
count	2.269900e+04	22699.000000	22699.000000	22699.000000	
mean	5.675849e+07	1.556236	1.642319	2.913313	
std	3.274493e+07	0.496838	1.285231	3.653171	
min	1.212700e+04	1.000000	0.000000	0.000000	
25%	2.852056e+07	1.000000	1.000000	0.990000	
50%	5.673150e+07	2.000000	1.000000	1.610000	
75%	8.537452e+07	2.000000	2.000000	3.060000	
max	1.134863e+08	2.000000	6.000000	33.960000	

	RatecodeID	PULocationID	DOLocationID	payment_type	fare_amount	\
count	22699.000000	22699.000000	22699.000000	22699.000000	22699.000000	
mean	1.043394	162.412353	161.527997	1.336887	13.026629	
std	0.708391	66.633373	70.139691	0.496211	13.243791	
min	1.000000	1.000000	1.000000	1.000000	-120.000000	
25%	1.000000	114.000000	112.000000	1.000000	6.500000	
50%	1.000000	162.000000	162.000000	1.000000	9.500000	
75%	1.000000	233.000000	233.000000	2.000000	14.500000	
max	99.000000	265.000000	265.000000	4.000000	999.990000	

	extra	mta_tax	tip_amount	tolls_amount	\
count	22699.000000	22699.000000	22699.000000	22699.000000	
mean	0.333275	0.497445	1.835781	0.312542	

std	0.463097	0.039465	2.800626	1.399212
min	-1.000000	-0.500000	0.000000	0.000000
25%	0.000000	0.500000	0.000000	0.000000
50%	0.000000	0.500000	1.350000	0.000000
75%	0.500000	0.500000	2.450000	0.000000
max	4.500000	0.500000	200.000000	19.100000

	improvement_surcharge	total_amount
count	22699.000000	22699.000000
mean	0.299551	16.310502
std	0.015673	16.097295
min	-0.300000	-120.300000
25%	0.300000	8.750000
50%	0.300000	11.800000
75%	0.300000	17.800000
max	0.300000	1200.290000

And info.

```
[16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Unnamed: 0                            22699 non-null  int64
1   VendorID                              22699 non-null  int64
2   tpep_pickup_datetime                  22699 non-null  object
3   tpep_dropoff_datetime                  22699 non-null  object
4   passenger_count                        22699 non-null  int64
5   trip_distance                          22699 non-null  float64
6   RatecodeID                            22699 non-null  int64
7   store_and_fwd_flag                    22699 non-null  object
8   PULocationID                          22699 non-null  int64
9   DOLocationID                          22699 non-null  int64
10  payment_type                           22699 non-null  int64
11  fare_amount                            22699 non-null  float64
12  extra                                  22699 non-null  float64
13  mta_tax                                22699 non-null  float64
14  tip_amount                             22699 non-null  float64
15  tolls_amount                           22699 non-null  float64
16  improvement_surcharge                  22699 non-null  float64
17  total_amount                           22699 non-null  float64
dtypes: float64(8), int64(7), object(3)
memory usage: 3.1+ MB
```

1.0.3 Task 2b. Select visualization type(s)

Now that you know which data columns you'll use, it is time to decide which data visualization makes the most sense for EDA of the TLC dataset. What type of data visualization(s) would be most helpful?

- Line graph
- Bar chart
- Box plot
- Histogram
- Heat map
- Scatter plot
- A geographic map

As you'll see below, a bar chart, box plot and scatter plot will be most helpful in your understanding of this data.

A box plot will be helpful to determine outliers and where the bulk of the data points reside in terms of trip_distance, duration, and total_amount

A scatter plot will be helpful to visualize the trends and patterns and outliers of critical variables, such as trip_distance and total_amount

A bar chart will help determine average number of trips per month, weekday, weekend, etc.

1.0.4 Task 3. Data visualization

You've assessed your data, and decided on which data variables are most applicable. It's time to plot your visualization(s)!

1.0.5 Boxplots

Perform a check for outliers on relevant columns such as trip distance and trip duration. Remember, some of the best ways to identify the presence of outliers in data are box plots and histograms.

Note: Remember to convert your date columns to datetime in order to derive total trip duration.

```
[19]: # Convert data columns to datetime
df["tpep_pickup_datetime"] = pd.to_datetime(df["tpep_pickup_datetime"])
df["tpep_dropoff_datetime"] = pd.to_datetime(df["tpep_dropoff_datetime"])
df.dtypes
```

```
[19]: Unnamed: 0                int64
VendorID                    int64
tpep_pickup_datetime        datetime64[ns]
tpep_dropoff_datetime        datetime64[ns]
passenger_count              int64
trip_distance                float64
RatecodeID                  int64
store_and_fwd_flag           object
PULocationID                 int64
DOLocationID                 int64
```

```

payment_type          int64
fare_amount           float64
extra                  float64
mta_tax                float64
tip_amount             float64
tolls_amount           float64
improvement_surcharge float64
total_amount           float64
dtype: object

```

trip distance

```

[21]: # Create box plot of trip_distance

# Define readable_numbers function
def readable_numbers(x):
    return f"{x:,.1f}" # Format with commas and 1 decimal

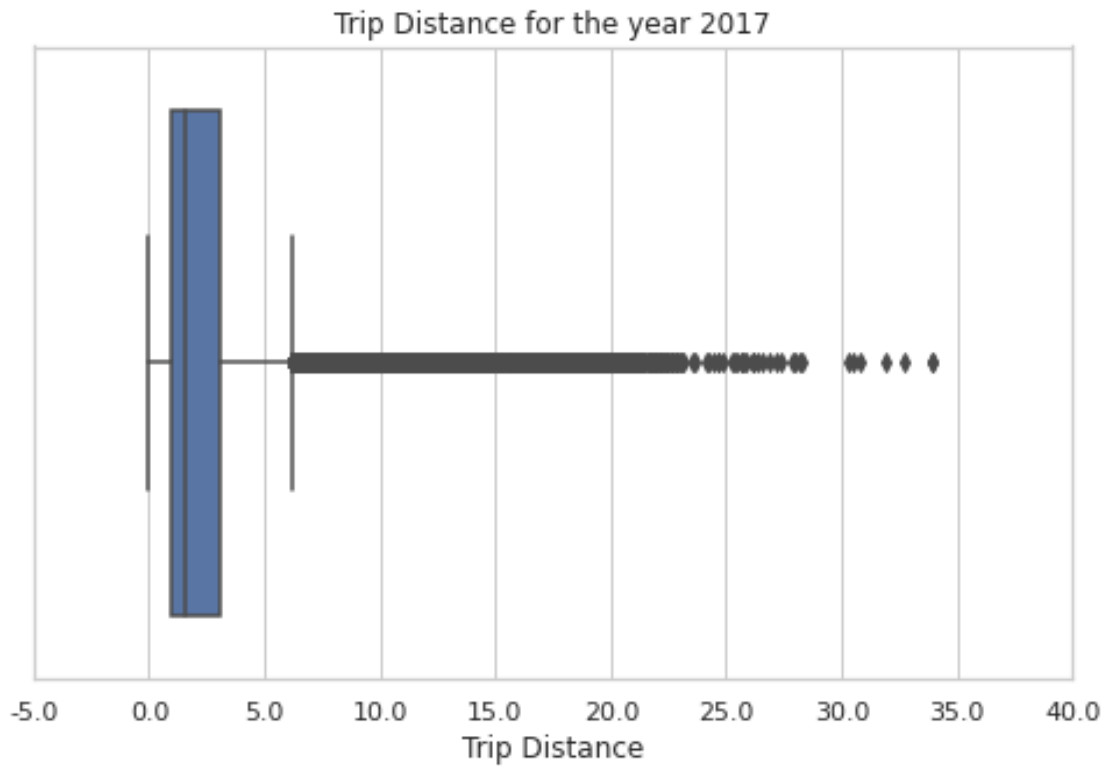
sns.set(style="whitegrid") # Set seaborn style

plt.figure(figsize=(8, 5)) # Set figure size
box = sns.boxplot(x=df['trip_distance']) # Create boxplot

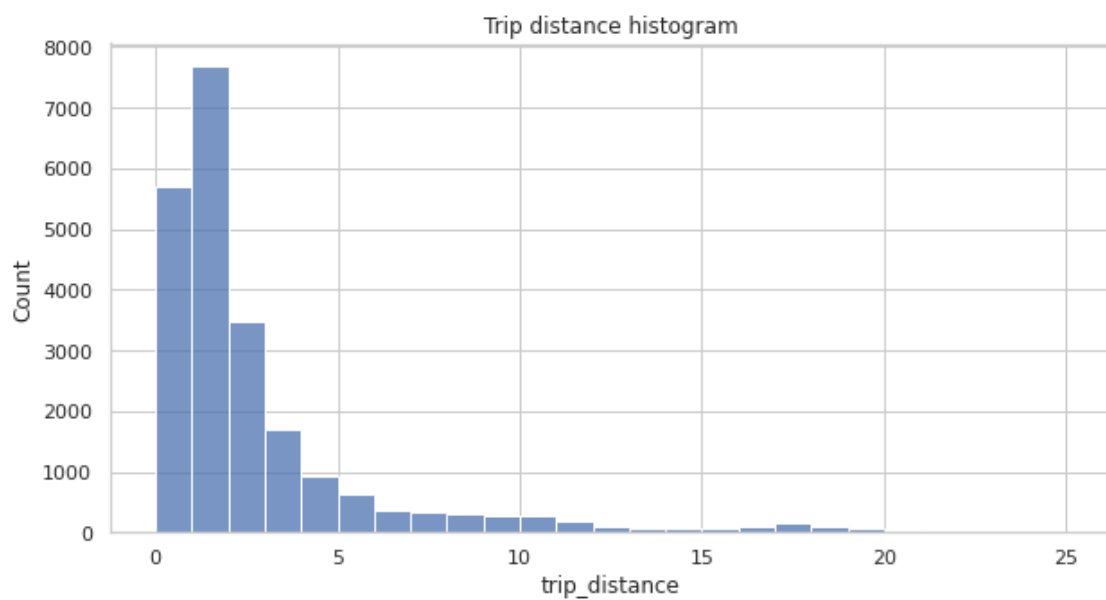
g = plt.gca() # Get current axes
plt.xticks(g.get_xticks(), [readable_numbers(x) for x in g.get_xticks()]) #
    ↪Format x-ticks
plt.xlabel('Trip Distance')
plt.title('Trip Distance for the year 2017')

plt.show() # Display plot

```



```
[31]: # Create histogram of trip_distance
plt.figure(figsize=(10,5))
sns.histplot(df['trip_distance'], bins=range(0,26,1))
plt.title('Trip distance histogram');
```



total amount

```
[25]: # Create box plot of total_amount

# Define readable_numbers function
def readable_numbers(x):
    return f"{x:,.1f}" # Format with commas and 1 decimal

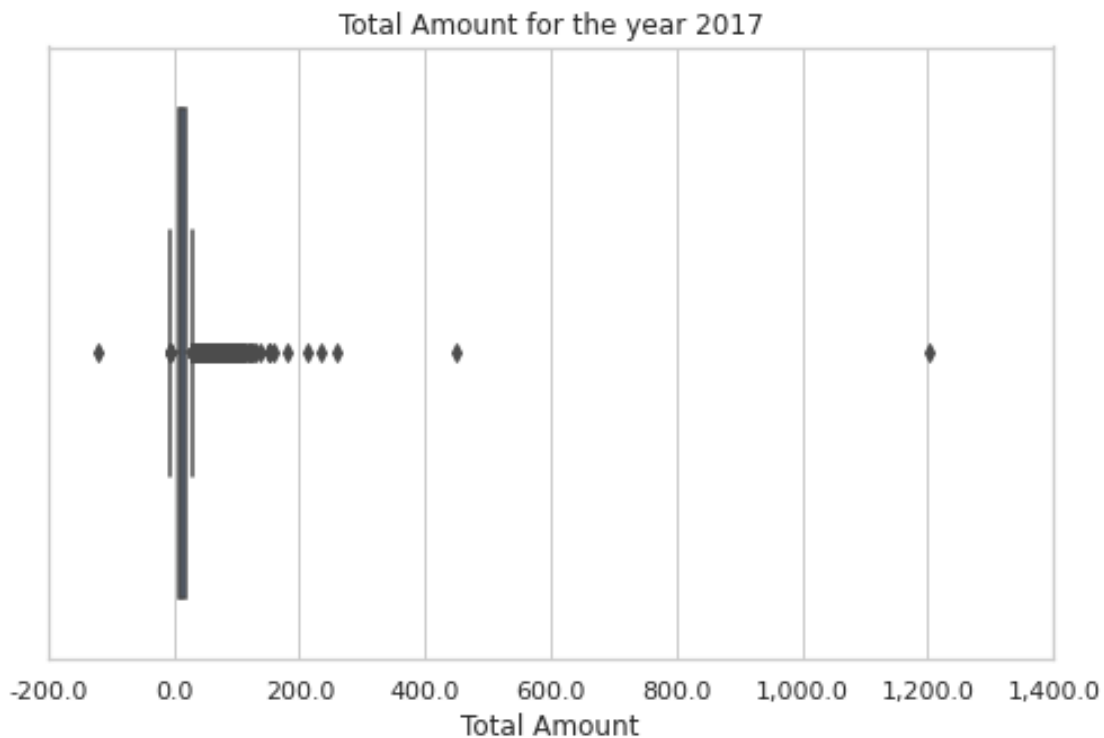
sns.set(style="whitegrid") # Set seaborn style

plt.figure(figsize=(8, 5)) # Set figure size
box = sns.boxplot(x=df['total_amount']) # Create boxplot

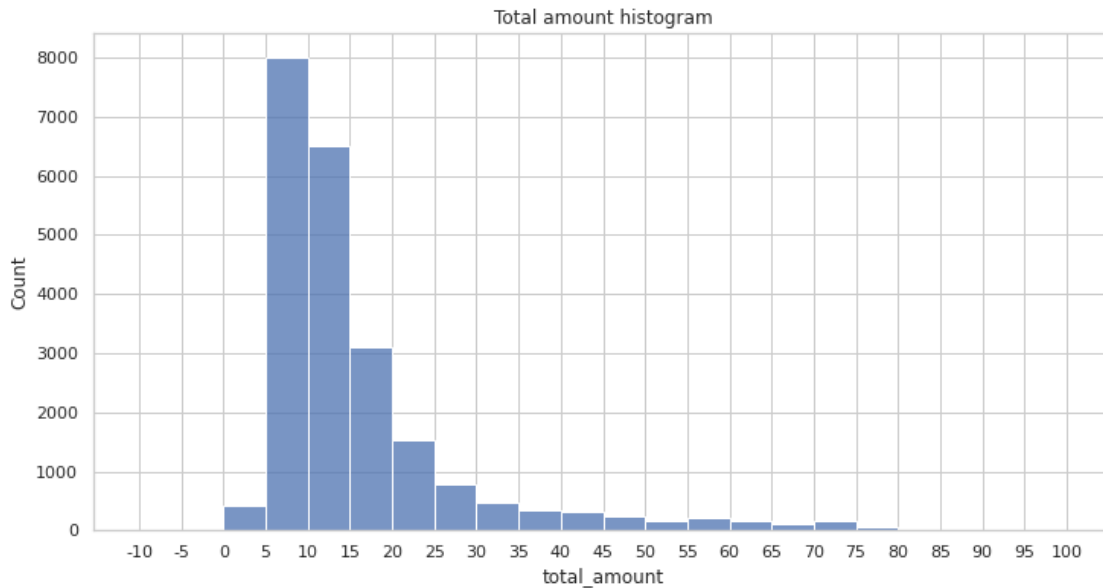
g = plt.gca() # Get current axes
plt.xticks(g.get_xticks(), [readable_numbers(x) for x in g.get_xticks()]) # ↵
    ↪Format x-ticks

plt.xlabel('Total Amount')
plt.title('Total Amount for the year 2017')

plt.show() # Display plot
```




```
[32]: # Create histogram of total_amount
plt.figure(figsize=(12,6))
ax = sns.histplot(df['total_amount'], bins=range(-10,101,5))
ax.set_xticks(range(-10,101,5))
ax.set_xticklabels(range(-10,101,5))
plt.title('Total amount histogram');
```



tip amount

```
[27]: # Create box plot of tip_amount

# Define readable_numbers function
def readable_numbers(x):
    return f"{x:,.1f}" # Format with commas and 1 decimal

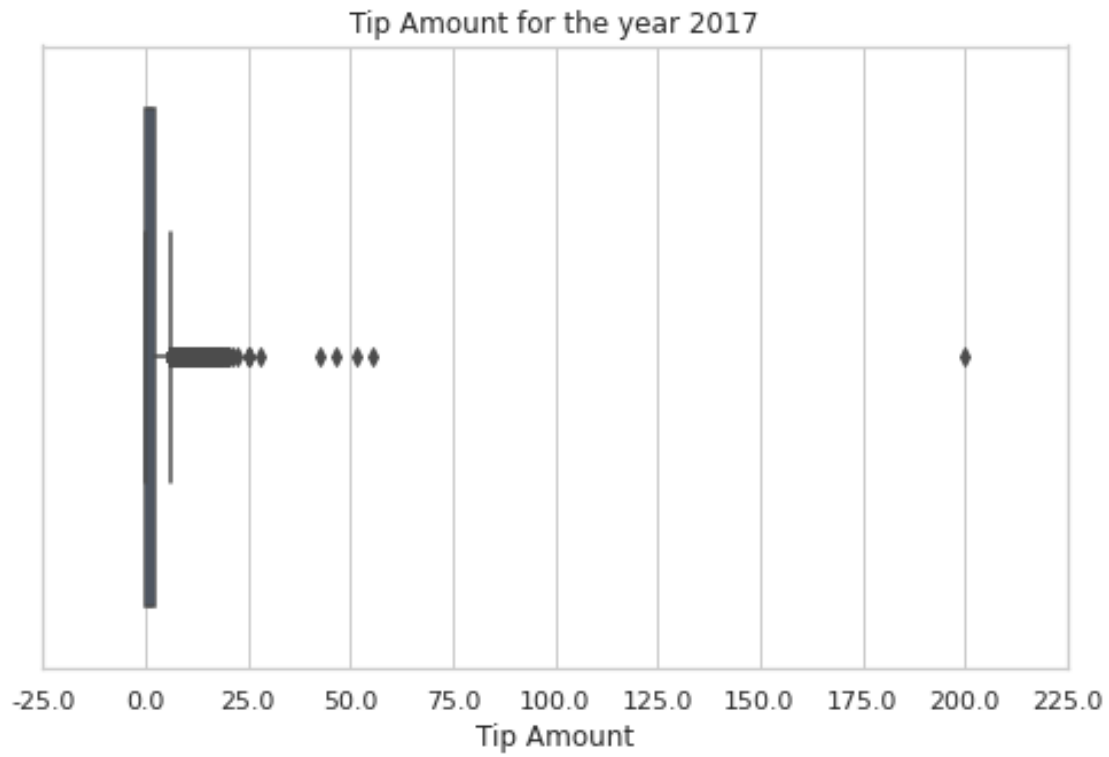
sns.set(style="whitegrid") # Set seaborn style

plt.figure(figsize=(8, 5)) # Set figure size
box = sns.boxplot(x=df['tip_amount']) # Create boxplot

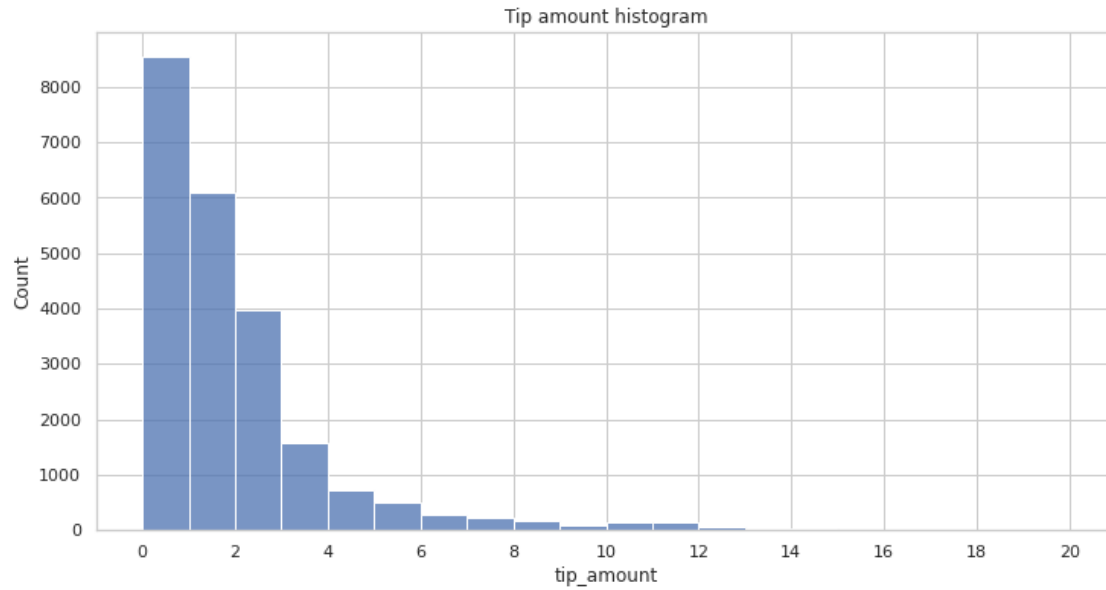
g = plt.gca() # Get current axes
plt.xticks(g.get_xticks(), [readable_numbers(x) for x in g.get_xticks()]) #_
    ↪Format x-ticks

plt.xlabel('Tip Amount')
plt.title('Tip Amount for the year 2017')

plt.show() # Display plot
```




```
[33]: # Create histogram of tip_amount
plt.figure(figsize=(12,6))
ax = sns.histplot(df['tip_amount'], bins=range(0,21,1))
ax.set_xticks(range(0,21,2))
ax.set_xticklabels(range(0,21,2))
plt.title('Tip amount histogram');
```




tip_amount by vendor

```
[34]: # Create histogram of tip_amount by vendor
plt.figure(figsize=(12,7))
ax = sns.histplot(data=df, x='tip_amount', bins=range(0,21,1),
                  hue='VendorID',
                  multiple='stack',
                  palette='pastel')
ax.set_xticks(range(0,21,1))
ax.set_xticklabels(range(0,21,1))
plt.title('Tip amount by vendor histogram');
```



Next, zoom in on the upper end of the range of tips to check whether vendor one gets noticeably more of the most generous tips.

```
[35]: # Create histogram of tip_amount by vendor for tips > $10
tips_over_ten = df[df['tip_amount'] > 10]
plt.figure(figsize=(12,7))
ax = sns.histplot(data=tips_over_ten, x='tip_amount', bins=range(10,21,1),
                  hue='VendorID',
                  multiple='stack',
                  palette='pastel')
ax.set_xticks(range(10,21,1))
ax.set_xticklabels(range(10,21,1))
plt.title('Tip amount by vendor histogram');
```



Mean tips by passenger count

Examine the unique values in the `passenger_count` column.

```
[36]: df['passenger_count'].value_counts()
```

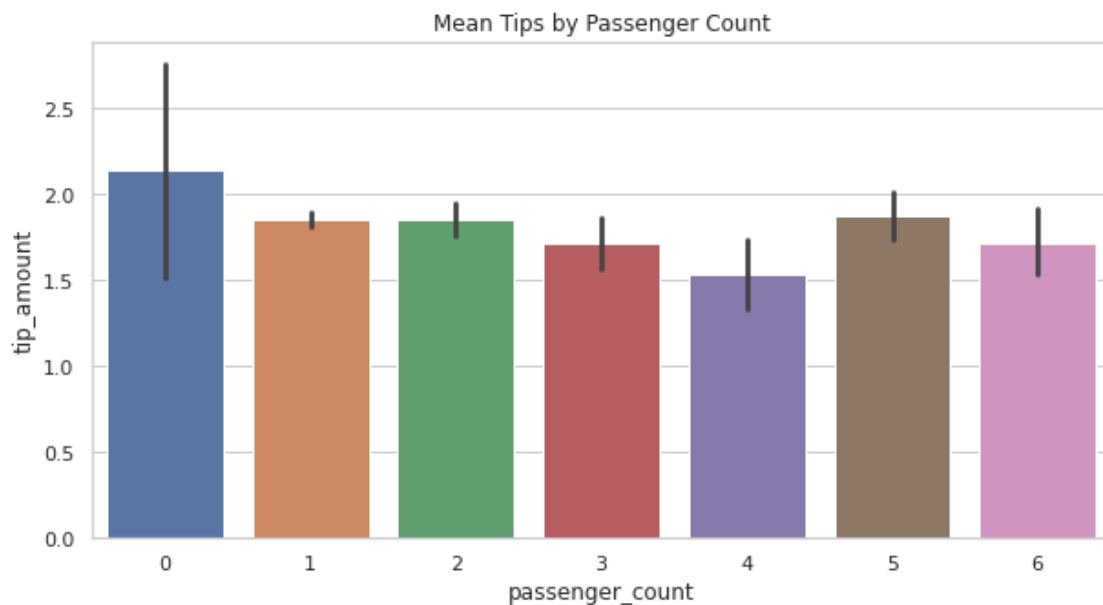
```
[36]: 1    16117
      2     3305
      5     1143
      3     953
      6     693
      4     455
      0       33
      Name: passenger_count, dtype: int64
```

```
[37]: # Calculate mean tips by passenger_count
mean_tips_by_passenger_count = df.groupby(['passenger_count']).
    .mean()['tip_amount']
mean_tips_by_passenger_count
```

```
[37]:          tip_amount
passenger_count
0          2.135758
1          1.848920
2          1.856378
3          1.716768
```

```
4          1.530264
5          1.873185
6          1.720260
```

```
[39]: # Create bar plot for mean tips by passenger count
plt.figure(figsize=(10, 5))
sns.barplot(x='passenger_count', y='tip_amount', data=df)
plt.title("Mean Tips by Passenger Count")
# Show the plot.
plt.show()
```



Create month and day columns

```
[41]: # Create a month column
df['month'] = df['tpep_pickup_datetime'].dt.month_name()
# Create a day column
df['day'] = df['tpep_pickup_datetime'].dt.day_name()
df.head()
```

```
[41]: Unnamed: 0  VendorID  tpep_pickup_datetime  tpep_dropoff_datetime  \
0    24870114         2  2017-03-25 08:55:43    2017-03-25 09:09:47
1    35634249         1  2017-04-11 14:53:28    2017-04-11 15:19:58
2    106203690         1  2017-12-15 07:26:56    2017-12-15 07:34:08
3    38942136         2  2017-05-07 13:17:59    2017-05-07 13:48:14
4    30841670         2  2017-04-15 23:32:20    2017-04-15 23:49:03

passenger_count  trip_distance  RatecodeID  store_and_fwd_flag  \
```

0	6	3.34	1	N
1	1	1.80	1	N
2	1	1.00	1	N
3	1	3.70	1	N
4	1	4.37	1	N

	PULocationID	DOLocationID	payment_type	fare_amount	extra	mta_tax	\
0	100	231	1	13.0	0.0	0.5	
1	186	43	1	16.0	0.0	0.5	
2	262	236	1	6.5	0.0	0.5	
3	188	97	1	20.5	0.0	0.5	
4	4	112	2	16.5	0.5	0.5	

	tip_amount	tolls_amount	improvement_surcharge	total_amount	month	\
0	2.76	0.0	0.3	16.56	March	
1	4.00	0.0	0.3	20.80	April	
2	1.45	0.0	0.3	8.75	December	
3	6.39	0.0	0.3	27.69	May	
4	0.00	0.0	0.3	17.80	April	

	day
0	Saturday
1	Tuesday
2	Friday
3	Sunday
4	Saturday

Plot total ride count by month

Begin by calculating total ride count by month.

```
[42]: # Get total number of rides for each month
monthly_rides = df['month'].value_counts()
monthly_rides
```

```
[42]: March      2049
      October    2027
      April      2019
      May        2013
      January    1997
      June       1964
      December   1863
      November   1843
      February   1769
      September  1734
      August     1724
      July       1697
      Name: month, dtype: int64
```

Reorder the results to put the months in calendar order.

```
[43]: # Reorder the monthly ride list so months go in order
month_order = ['January', 'February', 'March', 'April', 'May', 'June', 'July',
               'August', 'September', 'October', 'November', 'December']


monthly_rides = monthly_rides.reindex(index=month_order)
monthly_rides
```

```
[43]: January      1997
      February    1769
      March       2049
      April       2019
      May         2013
      June        1964
      July        1697
      August      1724
      September   1734
      October     2027
      November    1843
      December    1863
      Name: month, dtype: int64
```

```
[44]: # Show the index
monthly_rides.index
```

```
[44]: Index(['January', 'February', 'March', 'April', 'May', 'June', 'July',
        'August', 'September', 'October', 'November', 'December'],
        dtype='object')
```

```
[45]: # Create a bar plot of total rides per month
plt.figure(figsize=(12,7))
ax = sns.barplot(x=monthly_rides.index, y=monthly_rides)
ax.set_xticklabels(month_order)
plt.title('Ride count by month', fontsize=16);
```

Plot total ride count by day

Repeat the above process, but now calculate the total rides by day of the week.

```
[46]: # Repeat the above process, this time for rides by day
daily_rides = df['day'].value_counts()

day_order = ['Sunday', 'Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday']


daily_rides = daily_rides.reindex(index=day_order)

daily_rides
```

```
[46]: Sunday      2998
Monday      2931
Tuesday     3198
Wednesday   3390
Thursday    3402
Friday      3413
Saturday    3367
Name: day, dtype: int64
```

```
[47]: # Create bar plot for ride count by day
plt.figure(figsize=(12,7))
ax = sns.barplot(x=daily_rides.index, y=daily_rides)
ax.set_xticklabels(day_order)
```

```
plt.title('Ride count by day', fontsize=16);
```



Plot total revenue by day of the week

Repeat the above process, but now calculate the total revenue by day of the week.

```
[52]: # Repeat the process, this time for total revenue by day
daily_revenue = df.groupby("day")[["total_amount"]].sum()
print(daily_revenue)
```

```

total_amount
day
Friday      55818.74
Monday      49574.37
Saturday    51195.40
Sunday      48624.06
Thursday    57181.91
Tuesday     52527.14
Wednesday   55310.47

```

```
[56]: # Create bar plot of total revenue by day

# Reindex to match day order
daily_revenue = daily_revenue.reindex(index=day_order)


# Convert daily_revenue to 1D Series
daily_revenue_series = daily_revenue["total_amount"] # OR use .squeeze()
```

```

# Plot
plt.figure(figsize=(12, 7))
ax = sns.barplot(x=daily_revenue_series.index, y=daily_revenue_series.values)
# FIX APPLIED
ax.set_xticklabels(day_order)
plt.title("Total Revenue by Day", fontsize=16)
plt.xlabel("Day of the Week")
plt.ylabel("Total Revenue")

plt.show()

```



Plot total revenue by month

```

[57]: # Repeat the process, this time for total revenue by month
monthly_revenue = df.groupby("month")[["total_amount"]].sum()
print(monthly_revenue)

```

month	total_amount
April	32012.54
August	27759.56
December	31261.57
February	28937.89
January	31735.25
July	26617.64

June	32920.52
March	33085.89
May	33828.58
November	30800.44
October	33065.83
September	28206.38


```
[59]: # Create a bar plot of total revenue by month

# Reindex to match day order
monthly_revenue = monthly_revenue.reindex(index=month_order)

# Convert daily_revenue to 1D Series
monthly_revenue_series = monthly_revenue["total_amount"] # OR use .squeeze()

# Plot
plt.figure(figsize=(12, 7))
ax = sns.barplot(x=monthly_revenue_series.index, y=monthly_revenue_series.
    ↪values) # FIX APPLIED
ax.set_xticklabels(month_order)
plt.title("Total Revenue by Month", fontsize=16)
plt.xlabel("Months of the Year")
plt.ylabel("Total Revenue")

plt.show()
```



Plot mean trip distance by drop-off location

```
[60]: # Get number of unique drop-off location IDs
df['DOLocationID'].nunique()
```

```
[60]: 216
```


```
[61]: # Calculate the mean trip distance for each drop-off location
distance_by_dropoff = df.groupby('DOLocationID').mean()[['trip_distance']]

# Sort the results in descending order by mean trip distance
distance_by_dropoff = distance_by_dropoff.sort_values(by='trip_distance')
distance_by_dropoff
```

```
[61]:          trip_distance
DOLocationID
207          1.200000
193          1.390556
237          1.555494
234          1.727806
137          1.818852
...
51          17.310000
11          17.945000
210         20.500000
29          21.650000
23          24.275000
```

```
[216 rows x 1 columns]
```

```
[62]: # Create a bar plot of mean trip distances by drop-off location in ascending
      ↪order by distance
plt.figure(figsize=(14,6))
ax = sns.barplot(x=distance_by_dropoff.index,
                 y=distance_by_dropoff['trip_distance'],
                 order=distance_by_dropoff.index)
ax.set_xticklabels([])
ax.set_xticks([])
plt.title('Mean trip distance by drop-off location', fontsize=16);
```



You've now completed professional data visualizations according to a business need. Well done!